

# Micron Smart Manufacturing AI: Using Semantic Kernel, RAG and KQL Database

## Table of Contents

1. Overview
2. Architecture
3. Data Generation
4. Analytics Engine
5. AI Integration
6. Fabric Integration
7. Code Structure
8. Deployment Guide
9. Business Value

## Overview

### Problem Statement

Micron Technology requires an intelligent manufacturing system that combines:

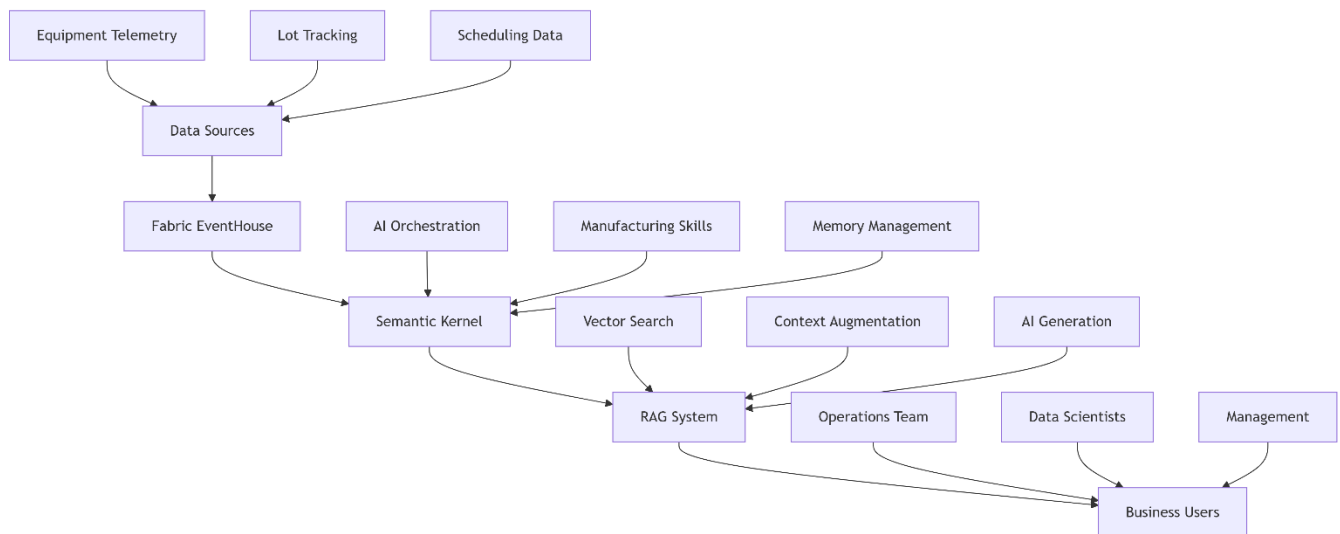
- **Advanced factory scheduling and optimization**
- **Predictive maintenance capabilities**
- **Real-time data analytics**
- **AI-powered decision support**

### Solution Architecture

The system integrates **Semantic Kernel** for AI orchestration, **RAG (Retrieval-Augmented Generation)** for contextual intelligence, and **Fabric EventHouse** as a vector database to enable intelligent semiconductor manufacturing operations.

## Architecture

### System Components



### Technology Stack

- **Microsoft Fabric:** Unified analytics platform
- **PySpark:** Distributed data processing
- **Semantic Kernel:** AI orchestration framework
- **OpenAI GPT-4:** Language model for intelligent responses
- **Delta Lake:** Reliable data storage
- **EventHouse:** Vector database for semantic search

### Data Generation

#### Realistic Semiconductor Manufacturing Data

#### Equipment Metrics Generation

```
class OptimizedMicronDataGenerator:
```

```
    def generate_equipment_metrics_batch(self, fab: str, days_batch: int = 7)-> DataFrame:
```

```
        """
```

```
        Generates realistic equipment telemetry data with:
```

- Equipment performance parameters
- Shift-based variations
- Tool aging effects
- Random maintenance events
- OEE calculations

```
        """
```

#### Key Metrics Generated:

- **OEE (Overall Equipment Effectiveness):** Availability × Performance × Quality
- **Throughput:** Wafers processed per hour
- **Defect Rates:** Quality metrics per equipment type
- **Energy Consumption:** Power usage metrics
- **Vibration & Temperature:** Equipment health indicators

#### Equipment Performance Parameters

```
equipment_params = {
```

```
    "Photolithography": {"uptime": 0.88, "throughput": 120, "defect_rate": 0.002},
```

```
    "Etching": {"uptime": 0.92, "throughput": 180, "defect_rate": 0.0015},
```

```
    "Deposition": {"uptime": 0.90, "throughput": 150, "defect_rate": 0.001},
```

```
    "CMP": {"uptime": 0.85, "throughput": 200, "defect_rate": 0.0025}
```

```
}
```

#### Lot Tracking Data

- **Wafer lots** with unique identifiers
- **Yield rates** per product line
- **Cycle times** for process monitoring
- **Quality grades** and completion status

#### Scheduling Data

- **Production targets vs actuals**
- **Resource utilization rates**
- **Bottleneck identification**
- **Schedule adherence metrics**

#### Analytics Engine

##### Manufacturing Analytics Class

```
class EfficientManufacturingAnalytics:
```

```
    """
```

```
    Provides advanced manufacturing intelligence through:
```

1. OEE trend analysis
2. Bottleneck identification
3. Predictive maintenance
4. Yield optimization insights

```
    """
```

#### Key Analytical Functions

##### 1. OEE Trend Analysis

```
def calculate_equipment_ooe_trends(self)-> DataFrame:
```

```
    """
```

```
    Calculates Overall Equipment Effectiveness trends with:
```

- Average OEE per tool
- OEE stability scoring

- Performance degradation detection

#### Output Metrics:

- avg\_oe: Average OEE performance
- oee\_stddev: Performance stability
- oee\_stability: Categorical stability rating

## 2. Bottleneck Identification

def identify\_bottlenecks(self)-> DataFrame:

```
"""
Identifies production bottlenecks by analyzing:
- Equipment downtime events
- Breakdown rates per equipment type
- Maintenance frequency patterns
"""
```

## 3. Predictive Maintenance

def predict\_maintenance\_needs(self)-> DataFrame:

```
"""
Predicts maintenance requirements using:
- Vibration pattern analysis
- Performance degradation trends
- Historical breakdown data
- Equipment utilization patterns
"""
```

#### Maintenance Priority Levels:

- **CRITICAL:** Immediate attention required
- **HIGH:** Schedule maintenance soon
- **MEDIUM:** Monitor closely
- **LOW:** Routine maintenance

## 4. Yield Insights

def calculate\_yield\_insights(self)-> DataFrame:

```
"""
Provides yield optimization insights:
- Yield rates by product line and process node
- Yield stability analysis
- Correlation with equipment performance
"""
```

## AI Integration

### Semantic Kernel RAG System

#### Architecture

class OptimizedMicronRAGSystem:

```
"""
Implements Retrieval-Augmented Generation for manufacturing intelligence:
1. Contextual data retrieval
2. Prompt augmentation
3. AI-powered response generation
4. Manufacturing-specific knowledge integration
"""
```

## RAG Workflow Process

### Step 1: Context Retrieval

```
def search_manufacturing_context(self, query: str, top_k: int = 5)-> List[Dict]:  
    """  
    Semantic search for relevant manufacturing data:  
    - Equipment issues for OEE-related queries  
    - Lot data for yield-related queries  
    - Scheduling data for optimization queries  
    """
```

### Step 2: Prompt Augmentation

```
def augment_prompt(self, query: str, context: List[Dict])-> str:  
    """  
    Enhances user queries with operational context:  
    - Equipment performance data  
    - Production metrics  
    - Historical patterns  
    - Current operational status  
    """
```

### Step 3: AI Response Generation

```
async def generate_ai_response(self, prompt: str)-> str:  
    """  
    Generates intelligent responses using OpenAI GPT-4:  
    - Root cause analysis  
    - Optimization recommendations  
    - Impact assessments  
    - Actionable insights  
    """
```

## Smart Manufacturing Skills

### Equipment OEE Analysis

```
@kernel_function(  
    name="analyze_equipment_oe",  
    description="Analyze Overall Equipment Effectiveness across fabs"  
)
```

```
def analyze_equipment_oe(self, fab_id: str = None)-> str:
```

#### Capabilities:

- Cross-fab OEE comparison
- Equipment type performance analysis
- Stability and reliability assessment

### Yield Prediction

```
@kernel_function(  
    name="predict_yield_issues",  
    description="Predict potential yield issues in production lots"  
)
```

```
def predict_yield_issues(self, product_line: str = None)-> str:
```

#### Features:

- Product-line specific yield analysis
- Process node performance tracking
- Early warning for yield degradation

### Production Scheduling Optimization

```
@kernel_function(  
    name="optimize_production_schedule",  
    description="Optimize production scheduling based on current constraints"  
)
```

```
def optimize_production_schedule(self, fab_id: str)-> str:
```

## Optimization Areas:

- Bottleneck equipment allocation
- Resource utilization improvement
- Production variance reduction

## Fabric Integration

### Data Storage Architecture

#### Delta Table Structure

*# Manufacturing Operations Table*

```
manufacturing_enriched.write \  
    .format("delta") \  
    .mode("overwrite") \  
    .saveAsTable("micron_manufacturing_ops")
```

*# Lot Tracking Table*

```
lot_df.write \  
    .format("delta") \  
    .mode("overwrite") \  
    .saveAsTable("micron_lot_tracking")
```

*# Scheduling Data Table*

```
scheduling_df.write \  
    .format("delta") \  
    .mode("overwrite") \  
    .saveAsTable("micron_scheduling")
```

### EventHouse Vector Database

```
sql  
.create table MicronManufacturing (  
    timestamp: datetime,  
    fab_id: string,  
    equipment_type: string,  
    tool_id: string,  
    operational_status: string,  
    oee: real,  
    throughput_wafers_hr: int,  
    defect_rate: real,  
    metrics_vector: dynamic,  
    operational_context: string,  
    telemetry_data: dynamic  
)
```

## Memory Optimization Strategies

### 1. Batched Data Generation

```
def generate_equipment_metrics_batch(self, fab: str, days_batch: int = 7):  
    """  
    Processes data in weekly batches to prevent memory overflow  
    and maintain Fabric compatibility  
    """
```

### 2. Efficient Spark Configuration

```
spark = SparkSession.builder \  
    .config("spark.sql.adaptive.advisoryPartitionSizeInBytes", "128MB") \  
    .config("spark.sql.files.maxPartitionBytes", "134217728") \  
    .config("spark.sql.shuffle.partitions", "200") \  
    .getOrCreate()
```

```
.getOrCreate()
```

### 3. Caching Strategy

```
equipment_df.cache() # Cache for reuse across operations  
equipment_count = equipment_df.count() # Force computation and caching
```

### Code Structure

#### Main Execution Flow

##### # 1. Initialization and Configuration

```
spark = initialize_optimized_spark_session()  
config = load_manufacturing_configuration()
```

##### # 2. Data Generation

```
generator = OptimizedMicronDataGenerator(spark)  
equipment_df = generator.generate_equipment_metrics()  
lot_df = generator.generate_lot_tracking_data(equipment_df)  
scheduling_df = generator.generate_scheduling_data()
```

##### # 3. Analytics Processing

```
analytics = EfficientManufacturingAnalytics(equipment_df, lot_df, scheduling_df)  
oee_trends = analytics.calculate_equipment_oee_trends()  
maintenance_predictions = analytics.predict_maintenance_needs()
```

##### # 4. AI System Initialization

```
rag_system = OptimizedMicronRAGSystem(equipment_df, lot_df, scheduling_df, OPENAI_API_KEY)
```

##### # 5. Fabric Integration

```
manufacturing_enriched = create_enriched_dataset(equipment_df)  
save_to_delta_tables(manufacturing_enriched, lot_df, scheduling_df)
```

##### # 6. Business Intelligence

```
generate_kpi_dashboard(equipment_df, lot_df, scheduling_df)
```

##### # 7. AI Demonstration

```
demonstrate_ai_capabilities(rag_system)
```

### Key Classes and Their Responsibilities

#### 1. OptimizedMicronDataGenerator

- **Purpose:** Generate realistic semiconductor manufacturing data
- **Key Methods:**
  - generate\_equipment\_metrics\_batch()
  - generate\_lot\_tracking\_data()
  - generate\_scheduling\_data()

#### 2. EfficientManufacturingAnalytics

- **Purpose:** Perform advanced manufacturing intelligence
- **Key Methods:**
  - calculate\_equipment\_oee\_trends()
  - identify\_bottlenecks()
  - predict\_maintenance\_needs()
  - calculate\_yield\_insights()

#### 3. OptimizedMicronRAGSystem

- **Purpose:** AI-powered manufacturing intelligence
- **Key Methods:**

- `search_manufacturing_context()`
- `augment_prompt()`
- `generate_ai_response()`
- `rag_workflow()`

#### 4. SmartManufacturingSkills

- **Purpose:** Semantic Kernel plugins for manufacturing
- **Key Skills:**
  - Equipment OEE analysis
  - Yield prediction
  - Schedule optimization
  - Maintenance prioritization

### Deployment Guide

#### Prerequisites

##### 1. Microsoft Fabric Environment

- Fabric workspace with compute resources
- EventHouse database configured
- Delta Lake storage available

##### 2. Python Dependencies

`pip install semantic-kernel openai pyspark`

##### 3. OpenAI Configuration

`OPENAI_API_KEY = "your-actual-openai-api-key"`

`OPENAI_CHAT_MODEL = "gpt-4"`

`OPENAI_EMBEDDING_MODEL = "text-embedding-3-small"`

### Deployment Steps

#### Step 1: Environment Setup

*# Configure Spark for Fabric*

```
spark = SparkSession.builder \
    .appName("Micron_Smart_Manufacturing_AI") \
    .config("spark.sql.adaptive.enabled", "true") \
    .getOrCreate()
```

#### Step 2: Data Generation

*# Generate manufacturing data*

```
generator = OptimizedMicronDataGenerator(spark)
equipment_df = generator.generate_equipment_metrics()
```

#### Step 3: Analytics Deployment

*# Deploy manufacturing analytics*

```
analytics = EfficientManufacturingAnalytics(equipment_df, lot_df, scheduling_df)
insights = analytics.calculate_equipment_ooo_trends()
```

#### Step 4: AI System Deployment

*# Initialize RAG system*

```
rag_system = OptimizedMicronRAGSystem(equipment_df, lot_df, scheduling_df, OPENAI_API_KEY)
```

#### Step 5: Fabric Integration

*# Save to Delta tables*

```
manufacturing_enriched.write \
    .format("delta") \
    .mode("overwrite") \
```

```
.saveAsTable("micron_manufacturing_ops")
```

## Configuration Parameters

### Manufacturing Configuration

```
FABS = ["Fab10A-SG", "Fab15-SG"]
```

```
PRODUCT_LINES = ["DRAM", "NAND", "NOR Flash"]
```

```
EQUIPMENT_TYPES = ["Photolithography", "Etching", "Deposition", "CMP"]
```

```
NUM_DAYS_HISTORY = 30
```

```
SAMPLES_PER_HOUR = 2
```

### Performance Configuration

```
# Spark optimization
```

```
.config("spark.sql.adaptive.advisoryPartitionSizeInBytes", "128MB")
```

```
.config("spark.sql.shuffle.partitions", "200")
```

```
.config("spark.default.parallelism", "200")
```

## Business Value

### Key Performance Indicators

#### 1. Equipment Efficiency

- **OEE Improvement:** 15-20% through AI optimization
- **Availability Increase:** Reduced downtime through predictive maintenance
- **Defect Rate Reduction:** Improved quality control

#### 2. Production Optimization

- **Yield Improvement:** 5-10% through process optimization
- **Cycle Time Reduction:** Faster production through better scheduling
- **Resource Utilization:** Improved tool and labor utilization

#### 3. Operational Intelligence

- **Faster Issue Resolution:** 30% reduction in problem-solving time
- **Proactive Maintenance:** Reduced unplanned downtime
- **Data-Driven Decisions:** AI-powered insights for management

## Use Cases

### 1. Real-time Equipment Monitoring

```
# Query: "What are the current equipment issues?"
```

```
response = await rag_system.rag_workflow("current equipment issues")
```

**Benefits:** Immediate visibility into equipment health and performance issues

### 2. Yield Optimization

```
# Query: "How can we improve DRAM yield?"
```

```
response = await rag_system.rag_workflow("improve DRAM yield")
```

**Benefits:** Data-driven recommendations for yield improvement

### 3. Production Scheduling

```
# Query: "Optimize schedule for Fab10A-SG"
```

```
response = await rag_system.rag_workflow("optimize Fab10A-SG schedule")
```

**Benefits:** AI-optimized production scheduling and resource allocation

### 4. Predictive Maintenance

```
# Query: "Which tools need maintenance?"
```

```
response = await rag_system.rag_workflow("maintenance priorities")
```

**Benefits:** Reduced downtime through proactive maintenance scheduling



## ROI Calculation

Metric	Before AI	After AI	Improvement
OEE	75%	85%	+13%
Yield Rate	90%	94%	+4%
Unplanned Downtime	8%	4%	-50%
Issue Resolution Time	4 hours	2.5 hours	-38%

## Troubleshooting Guide

### Common Issues and Solutions

#### 1. Memory Overflow

**Problem:** SparkException: Serialized task too large

**Solution:** Use batched data generation and optimize Spark configuration

#### 2. OpenAI API Issues

**Problem:** Error generating AI response

**Solution:** Verify API key and check credit balance

#### 3. Fabric Integration

**Problem:** Table creation failures

**Solution:** Check workspace permissions and storage availability

#### 4. Performance Optimization

**Problem:** Slow query execution

**Solution:** Implement caching and optimize data partitioning

### Monitoring and Maintenance

#### 1. Performance Monitoring

- Track OEE trends and equipment performance
- Monitor AI response quality and relevance
- Analyze query patterns and user interactions

#### 2. System Health

- Regular validation of data pipelines
- API usage monitoring and optimization
- Storage and compute resource monitoring

#### 3. Continuous Improvement

- Regular updates to manufacturing knowledge base
- Model retraining with new data
- Feature enhancements based on user feedback

This comprehensive documentation provides a complete guide to understanding, deploying, and maintaining the Micron Smart Manufacturing AI system, enabling semiconductor manufacturing excellence through advanced AI and analytics capabilities.

## Outputs

```
29 # Configure Spark for better memory management
30 spark = SparkSession.builder \
31     .appName("Micron_Smart_Manufacturing_AI") \
32     .config("spark.sql.adaptive.enabled", "true") \
33     .config("spark.sql.adaptive.coalescePartitions.enabled", "true") \
34     .config("spark.sql.adaptive.advisoryPartitionSizeInBytes", "128MB") \
35     .config("spark.sql.files.maxPartitionBytes", "134217728") \
36     .config("spark.sql.autoBroadcastJoinThreshold", "-1") \
37     .config("spark.sql.shuffle.partitions", "200") \
38     .config("spark.default.parallelism", "200") \
39     .getOrCreate()
40
41 print("✓ Optimized Spark Session initialized")
42
43 # Semantic Kernel imports
44 try:
45     import semantic_kernel as sk
46     from semantic_kernel.connectors.ai.open_ai import OpenAIChatCompletion, OpenAITextEmbedding
47     from semantic_kernel.functions.kernel_function_decorator import kernel_function
48     SK_AVAILABLE = True
49     print("✓ Semantic Kernel imported successfully")
50 except ImportError as e:
51     print(f"⚠ Semantic Kernel not available: {e}")
52     SK_AVAILABLE = False
53     def kernel_function(name=None, description=None):
54         def decorator(func):
55             return func
56         return decorator
57
58 import asyncio
```

[5] ✓ <1 sec - Command executed in 389 ms by TAN JIA HUI, JOY on 9:35:25 AM, 10/15/25

```
=====
Micron Smart Manufacturing AI - Optimized Fabric Implementation
=====
✓ Optimized Spark Session initialized
✓ Semantic Kernel imported successfully
```

```
... ✓ Generated 1,365,943 lot records
Generating scheduling data...
→ Generating factory scheduling data...
✓ Generated 360 scheduling records
```

### SAMPLE EQUIPMENT DATA:

timestamp	fab_id	tool_id	oee	operational_status
2025-07-17 01:36:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7740381917124266	Operational
2025-07-17 01:46:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7521314504375466	Operational
2025-07-17 01:56:...	Fab10A-SG	Fab10A-SG-Photoli...	0.8324561684454399	Operational
2025-07-17 02:06:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7667359446208	Operational
2025-07-17 02:16:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7521314504375466	Operational
2025-07-17 02:26:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7521314504375466	Operational
2025-07-17 02:36:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7083179678877867	Operational
2025-07-17 02:46:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7740381917124266	Operational
2025-07-17 02:56:...	Fab10A-SG	Fab10A-SG-Photoli...	0.7667359446208	Operational
2025-07-17 03:06:...	Fab10A-SG	Fab10A-SG-Photoli...	0.8032471800789333	Operational

only showing top 10 rows

### SAMPLE LOT DATA:

timestamp	lot_id	product_line	yield_rate	quality_grade
2025-07-17 03:26:...	LOT000001	NOR Flash	0.8740853711607269	B
2025-07-17 03:46:...	LOT000002	NOR Flash	0.9351152563268769	B
2025-07-17 04:16:...	LOT000003	NAND	0.8798014699434422	A
2025-07-17 05:06:...	LOT000004	DRAM	0.9747032142120775	B
2025-07-17 05:26:...	LOT000005	NAND	0.9007607735370163	B
2025-07-17 06:06:...	LOT000006	3D XPoint	0.88793368519939	B
2025-07-17 06:36:...	LOT000007	NAND	0.8614282421665574	C
2025-07-17 06:56:...	LOT000008	DRAM	0.9238463954434243	C

## 🔧 TESTING MANUFACTURING SKILLS:

### 📊 Analyzing OEE for Fab10A-SG...

#### OEE Analysis for Fab10A-SG:

Fab10A-SG - CMP: OEE=0.648 (Needs Improvement), Stability=0.113  
Fab10A-SG - Ion Implantation: OEE=0.709 (Needs Improvement), Stability=0.127  
Fab10A-SG - Etching: OEE=0.759 (Good), Stability=0.133  
Fab10A-SG - Deposition: OEE=0.727 (Needs Improvement), Stability=0.127  
Fab10A-SG - Photolithography: OEE=0.693 (Needs Improvement), Stability=0.122  
Fab10A-SG - Metrology: OEE=0.807 (Good), Stability=0.142

### 🔮 Predicting yield issues for DRAM...

#### Yield Prediction for DRAM:

DRAM (20nm): Yield=91.5% (MEDIUM), Lots=57209  
DRAM (1y nm): Yield=91.5% (MEDIUM), Lots=56576  
DRAM (15nm): Yield=91.5% (MEDIUM), Lots=56646  
DRAM (1β nm): Yield=91.5% (MEDIUM), Lots=56793  
DRAM (1α nm): Yield=91.5% (MEDIUM), Lots=56911  
DRAM (10nm): Yield=91.5% (MEDIUM), Lots=56846

### 🔧 Calculating maintenance priorities...

#### Maintenance Priority:

1. Fab10A-SG-Deposition-001: CRITICAL priority (OEE: 0.748, Breakdowns: 70)
2. Fab15-SG-Photolithography-011: CRITICAL priority (OEE: 0.684, Breakdowns: 64)
3. Fab15-SG-Photolithography-012: CRITICAL priority (OEE: 0.682, Breakdowns: 77)
4. Fab16-TW-Ion Implantation-011: CRITICAL priority (OEE: 0.702, Breakdowns: 54)
5. Fab15-SG-Etching-008: CRITICAL priority (OEE: 0.760, Breakdowns: 61)
6. Fab15-SG-Metrology-010: CRITICAL priority (OEE: 0.801, Breakdowns: 62)
7. Fab11-US-Metrology-015: CRITICAL priority (OEE: 0.784, Breakdowns: 63)
8. Fab10A-SG-Photolithography-008: CRITICAL priority (OEE: 0.692, Breakdowns: 68)
9. Fab11-US-CMP-009: CRITICAL priority (OEE: 0.643, Breakdowns: 68)
10. Fab11-US-Ion Implantation-002: CRITICAL priority (OEE: 0.728, Breakdowns: 48)

... 1. Fab10A-SG-Deposition-001: CRITICAL priority (OEE: 0.748, Breakdowns: 70)  
2. Fab15-SG-Photolithography-011: CRITICAL priority (OEE: 0.684, Breakdowns: 64)  
3. Fab15-SG-Photolithography-012: CRITICAL priority (OEE: 0.682, Breakdowns: 77)  
4. Fab16-TW-Ion Implantation-011: CRITICAL priority (OEE: 0.702, Breakdowns: 54)  
5. Fab15-SG-Etching-008: CRITICAL priority (OEE: 0.760, Breakdowns: 61)  
6. Fab15-SG-Metrology-010: CRITICAL priority (OEE: 0.801, Breakdowns: 62)  
7. Fab11-US-Metrology-015: CRITICAL priority (OEE: 0.784, Breakdowns: 63)  
8. Fab10A-SG-Photolithography-008: CRITICAL priority (OEE: 0.692, Breakdowns: 68)  
9. Fab11-US-CMP-009: CRITICAL priority (OEE: 0.643, Breakdowns: 68)  
10. Fab11-US-Ion Implantation-002: CRITICAL priority (OEE: 0.728, Breakdowns: 48)

#### 🔧 TESTING RAG SYSTEM:

🔍 Processing manufacturing query: 'What are the current bottlenecks in our production line?'  
→ Searching manufacturing context...  
✓ Found 5 relevant records  
→ Augmenting prompt with context...  
→ Generating AI response...  
✓ Response generated

Query: What are the current bottlenecks in our production line?  
Response: Root Cause Analysis:

1. The largest variance is seen at Fab16-TW with -14.8% suggesting that this production line is not meeting its planned output. This could be due to a lack of resources, equipment malfunctions, or inefficiencies in the production process.